

# CAeSaR: unified Cluster-Assignment Scheduling and communication Reuse for clustered VLIW processors

Vasileios Porpodas<sup>†</sup> and Marcelo Cintra<sup>†\*</sup>

School of Informatics, University of Edinburgh<sup>†</sup>  
Intel Labs Braunschweig<sup>\*</sup>

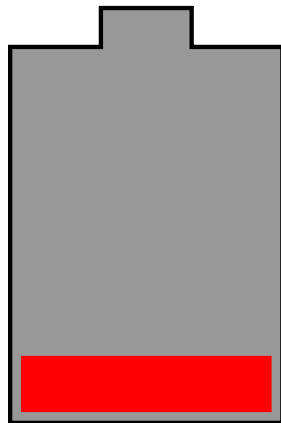
CASES 2013

# Energy efficiency, VLIWs and Scheduling

- Mobile era.

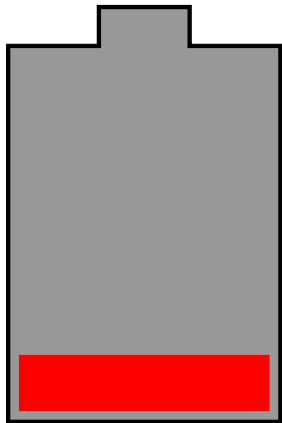
# Energy efficiency, VLIWs and Scheduling

- Mobile era.
- Energy becomes a major design constraint.



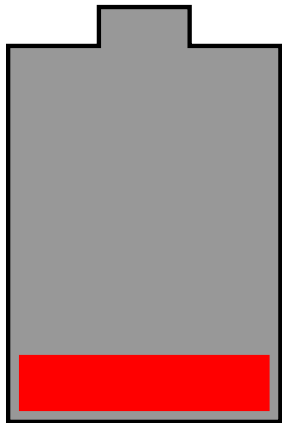
# Energy efficiency, VLIWs and Scheduling

- Mobile era.
- Energy becomes a major design constraint.
- Hardware Instruction scheduling consumes a lot of energy.



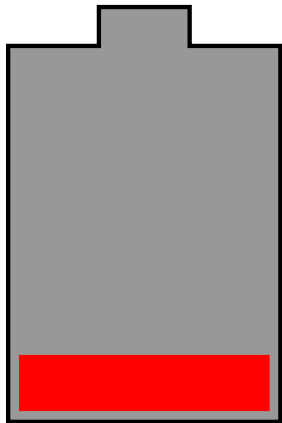
# Energy efficiency, VLIWs and Scheduling

- Mobile era.
- Energy becomes a major design constraint.
- Hardware Instruction scheduling consumes a lot of energy.
- VLIW processors: high-performance and statically scheduled.




# Energy efficiency, VLIWs and Scheduling

- Mobile era.
- Energy becomes a major design constraint.
- Hardware Instruction scheduling consumes a lot of energy.
- VLIW processors: high-performance and statically scheduled.
- Clustered VLIW = scalable VLIW.




# Clustered VLIW

- Scalable
- Energy efficient



# Clustered VLIW


- Scalable
- Energy efficient
- High frequency
- Inter-Cluster delay





# Clustered VLIW


- Scalable
- Energy efficient
- High frequency
- Inter-Cluster delay
- Relies on compiler
- Statically scheduled
- Explicit ILP



# Scheduling Overview

- Scheduler for monolithic VLIW

## Non-Clustered




**a. Compilation passes for non-clustered architectures.**

## Clustered Architecture

# Scheduling Overview


- Scheduler for monolithic VLIW
- 2-stage scheduling for clusters

## Non-Clustered



**a. Compilation passes for non-clustered architectures.**

## Clustered Architecture




**b. Decoupled Cluster Assignment and Scheduling.**

# Scheduling Overview


- Scheduler for monolithic VLIW
- 2-stage scheduling for clusters

## Non-Clustered



a. Compilation passes for non-clustered architectures.

## Clustered Architecture




b. Decoupled Cluster Assignment and Scheduling.

# Scheduling Overview


- Scheduler for monolithic VLIW
- 2-stage scheduling for clusters
- Unified scheduler (State-of-the-art)

## Non-Clustered



a. Compilation passes for non-clustered architectures.

## Clustered Architecture



b. Decoupled Cluster Assignment and Scheduling.




c. UAS: Unified cluster Assignment and Scheduling.

# Scheduling Overview


- Scheduler for monolithic VLIW
- 2-stage scheduling for clusters
- Unified scheduler (State-of-the-art)
- 2-stage ICC-reuse

## Non-Clustered



a. Compilation passes for non-clustered architectures.

## Clustered Architecture



b. Decoupled Cluster Assignment and Scheduling.



c. UAS: Unified cluster Assignment and Scheduling.




d. Decoupled UAS and ICC-Reuse.

# Scheduling Overview


- Scheduler for monolithic VLIW
- 2-stage scheduling for clusters
- Unified scheduler (State-of-the-art)
- 2-stage ICC-reuse

## Non-Clustered




a. Compilation passes for non-clustered architectures.

## Clustered Architecture



b. Decoupled Cluster Assignment and Scheduling.



c. UAS: Unified cluster Assignment and Scheduling.




d. Decoupled UAS and ICC-Reuse.

# Scheduling Overview


- Scheduler for monolithic VLIW
- 2-stage scheduling for clusters
- Unified scheduler (State-of-the-art)
- 2-stage ICC-reuse
- Unified ICC-reuse

## Non-Clustered




a. Compilation passes for non-clustered architectures.

## Clustered Architecture



b. Decoupled Cluster Assignment and Scheduling.



c. UAS: Unified cluster Assignment and Scheduling.



d. Decoupled UAS and ICC-Reuse.



e. CAeSaR: unified Cluster Assignment Scheduling & ICC-Reuse



# Outline


Introduction

CAeSaR Scheduling

Experimental Setup and Results

Conclusion


# CAeSaR Scheduler



- State-of-the-art (UAS)

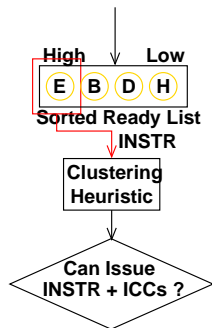
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR




# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR



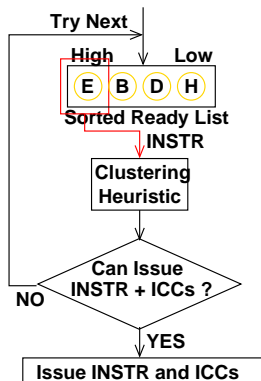
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR



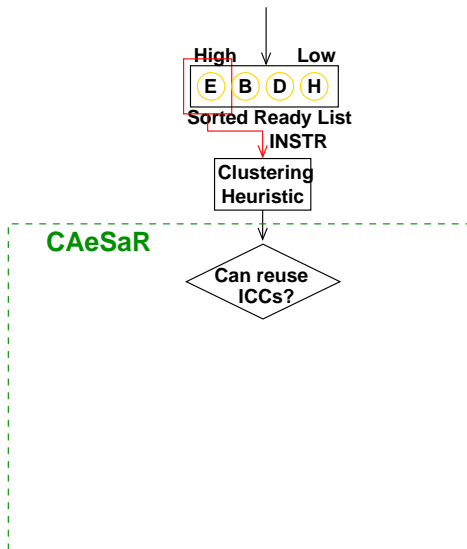
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE



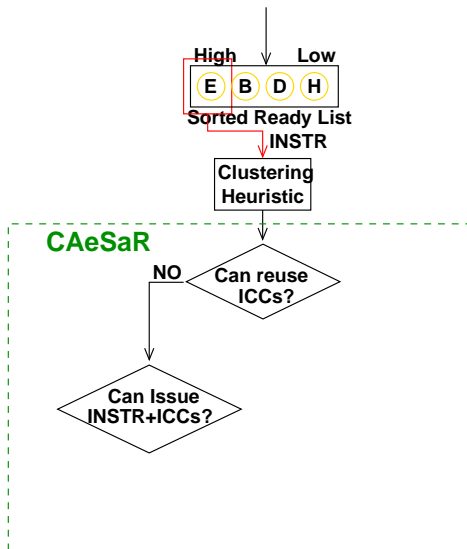
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE
- CAeSaR



# CAeSaR Scheduler

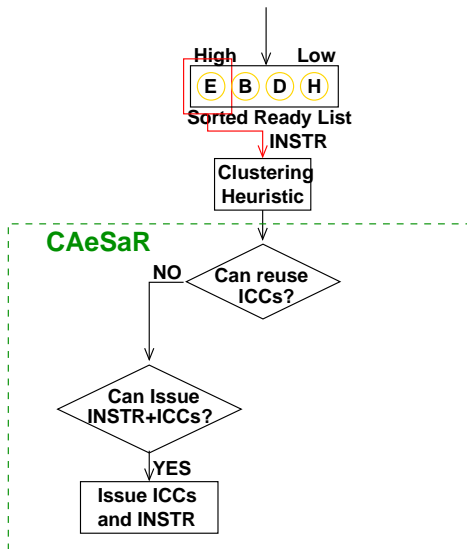
- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE
- CAeSaR





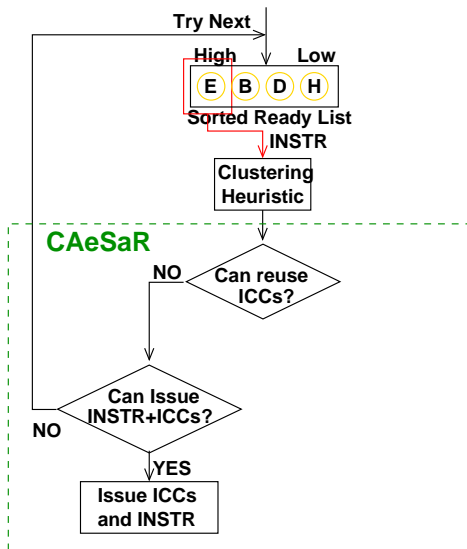
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE
- CAeSaR



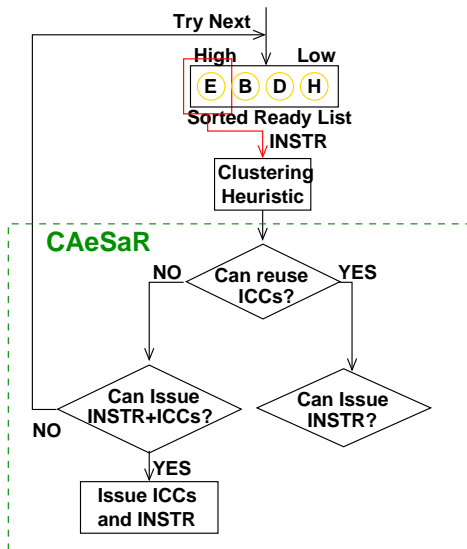
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE
- CAeSaR



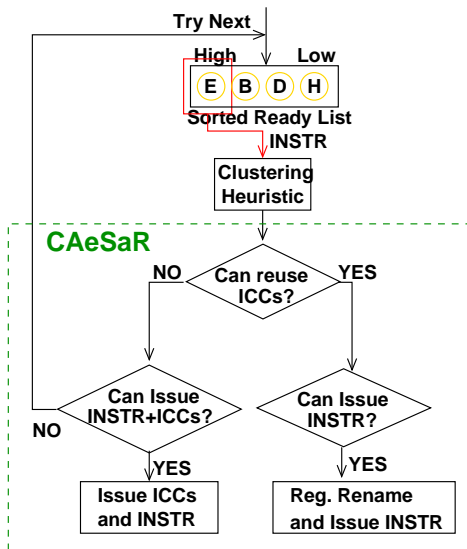
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE
- CAeSaR
- ICC-reuse



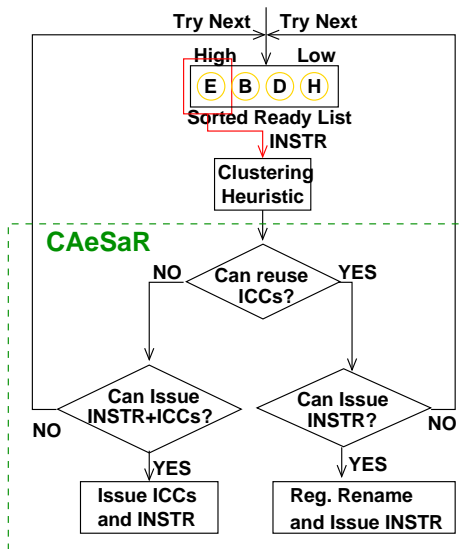
# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE
- CAeSaR
- ICC-reuse



# CAeSaR Scheduler

- State-of-the-art (UAS)
- High Priority INSTR
- If can issue, DONE
- CAeSaR
- ICC-reuse




# Example

## Monolithic

0	A	
1	B	F
2	C	G
3	D	H
4	I	J
5	E	

a. VLIW

## Clustered Architecture



f. Data Flow Graph (DFG)

# Example

- 2-step scheduler generates bad schedule

## Monolithic


0	A	
1	B	F
2	C	G
3	D	H
4	I	J
5	E	

a. VLIW

## Clustered Architecture

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	ICC
4	H	G
5		ICC
6		I
7		ICC
8		J
9	E	

b. Decoupled



f. Data Flow Graph (DFG)

## Example

- 2-step scheduler generates bad schedule
- UAS improves it by being ICC-aware

### Monolithic

0	A	
1	B	F
2	C	G
3	D	H
4	I	J
5	E	

a. VLIW


### Clustered Architecture

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	ICC
4	H	G
5		ICC
6		I
7		ICC
8		J
9	E	

b. Decoupled

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	ICC
4	G	H
5	I	
6	J	
7	E	

c. UAS



f. Data Flow Graph (DFG)



## Example

- 2-step scheduler generates bad schedule
- UAS improves it by being ICC-aware
- UAS with ICC-reuse has fewer ICCs

### Monolithic

0	A	
1	B	F
2	C	G
3	D	H
4	I	J
5	E	

a. VLIW

### Clustered Architecture

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	ICC
4	H	G
5		ICC
6		I
7		ICC
8		J
9	E	


b. Decoupled

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	ICC
4	G	H
5	I	
6	J	
7	E	

c. UAS

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	
4	G	H
5	I	
6	J	
7	E	

d. UAS+ICC-reuse



f. Data Flow Graph (DFG)

## Example

- 2-step scheduler generates bad schedule
- UAS improves it by being ICC-aware
- UAS with ICC-reuse has fewer ICCs
- CAeSaR solves phase-ordering

### Monolithic

0	A	
1	B	F
2	C	G
3	D	H
4	I	J
5	E	

a. VLIW

### Clustered Architecture

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	ICC
4	H	G
5		ICC
6		I
7		ICC
8		J
9	E	

b. Decoupled

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	ICC
4	G	H
5	I	
6	J	
7	E	


c. UAS

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	
4	G	H
5	I	
6	J	
7	E	

d. UAS+ICC-reuse

	CL0	CL1
0	A	
1	B	ICC
2	C	F
3	D	G
4	H	I
5	J	
6	E	


e. CAeSaR



f. Data Flow Graph (DFG)

# Register File Coherence


- Reused value must be correct




**a. Before ICCs**

# Register File Coherence

- Reused value must be correct




a. Before ICCs




b. Wrong: No Coherence

# Register File Coherence

- Reused value must be correct




a. Before ICCs




b. Wrong: No Coherence

# Register File Coherence


- Reused value must be correct
- CAeSaR follows an approach similar to write-invalidate cache coherent protocols



a. Before ICCs




b. Wrong: No Coherence




c. CAeSaR Register Coherence

# Register File Coherence


- Reused value must be correct
- CAeSaR follows an approach similar to write-invalidate cache coherent protocols



a. Before ICCs




b. Wrong: No Coherence




c. CAeSaR Register Coherence

# Register File Coherence


- Reused value must be correct
- CAeSaR follows an approach similar to write-invalidate cache coherent protocols



a. Before ICCs



b. Wrong: No Coherence



c. CAeSaR Register Coherence



# Outline

Introduction


CAeSaR Scheduling

Experimental Setup and Results

Conclusion


# Experimental Setup

- Compiler: GCC-4.5.0, Modified Haifa-Scheduler



## Experimental Setup


- Compiler: GCC-4.5.0, Modified Haifa-Scheduler



- Architecture
  - IA64-based 4 issue clustered VLIW
  - 2,4 clusters
  - 1 cycle Inter-Cluster Latency

## Experimental Setup


- Compiler: GCC-4.5.0, Modified Haifa-Scheduler



- Architecture
  - IA64-based 4 issue clustered VLIW
  - 2,4 clusters
  - 1 cycle Inter-Cluster Latency
- Benchmarks: SPEC CINT2000 and MediabenchII Video


## Results: 4-cluster 4-issue, ICC ratio

- ICCs are a significant bottleneck




## Results: 4-cluster 4-issue, ICC ratio

- ICCs are a significant bottleneck
- 1 ICC in 5 non-ICCs




## Results: 4-cluster 4-issue, Sched cycles

- CAeSaR generates 13.8% more compact schedules on average



## Results: 4-cluster 4-issue, Sched cycles


- CAeSaR generates 13.8% more compact schedules on average
- ICC-reuse same performance as UAS (expected)






## Results: 4-cluster 4-issue, ICCs

- CAeSaR reduces the count of ICCs




## Results: 4-cluster 4-issue, ICCs

- CAeSaR reduces the count of ICCs
- CAeSaR generates fewer ICCs than ICC-reuse (proof of phase-ordering problem)




## Results: 4-cluster 4-issue, Instr Distribution

- Better resource utilization
- Less communication bottlenecks
- Potential for more ILP




## Results: 2-cluster 4-issue, ICC ratio

- ICCs are still noticeable




## Results: 2-cluster 4-issue, ICC ratio

- ICCs are still noticable
- 1 ICC in 10 non-ICCs




## Results: 2-cluster 4-issue, Sched cycles

- CAeSaR generates 8.4% more compact schedules on average




## Results: 2-cluster 4-issue, Sched cycles

- CAeSaR generates 8.4% more compact schedules on average
- ICC-reuse same performance as UAS (expected)



## Results: 2-cluster 4-issue, ICCs


- CAeSaR reduces the count of ICCs





## Results: 2-cluster 4-issue, ICCs

- CAeSaR reduces the count of ICCs
- CAeSaR generates fewer ICCs than ICC-reuse (proof of phase-ordering problem)



## Conclusion

Proposed CAeSaR Scheduling, a scheduler for clustered VLIWs that:

- Eliminates redundant Inter-Cluster copies
- Solves phase-ordering problem between Scheduling and ICC-reuse
- Generates more compact schedules compared to state-of-the-art

# CAeSaR: unified Cluster-Assignment Scheduling and communication Reuse for clustered VLIW processors

Vasileios Porpodas<sup>†</sup> and Marcelo Cintra<sup>†\*</sup>

School of Informatics, University of Edinburgh<sup>†</sup>  
Intel Labs Braunschweig<sup>\*</sup>

CASES 2013