

PSLP: Padded SLP Automatic Vectorization

Vasileios Porpodas[†], Alberto Magni[‡]
and Timothy M. Jones[†]

University of Cambridge[†]
University of Edinburgh[‡]

CGO 2015



THE UNIVERSITY of EDINBURGH
informatics



UNIVERSITY OF
CAMBRIDGE

SLP Straight-Line Code Vectorizer

- Superword Level Parallelism [Larsen PLDI'00]

SLP Straight-Line Code Vectorizer

- Superword Level Parallelism [Larsen PLDI'00]
- State-of-the-art straight-line code vectorizer

SLP Straight-Line Code Vectorizer

- Superword Level Parallelism [Larsen PLDI'00]
- State-of-the-art straight-line code vectorizer
- Implemented in most compilers (including GCC and LLVM)

SLP Straight-Line Code Vectorizer

- Superword Level Parallelism [Larsen PLDI'00]
- State-of-the-art straight-line code vectorizer
- Implemented in most compilers (including GCC and LLVM)
- In theory it should be a superset of loop-vectorizer

SLP Straight-Line Code Vectorizer

- Superword Level Parallelism [Larsen PLDI'00]
- State-of-the-art straight-line code vectorizer
- Implemented in most compilers (including GCC and LLVM)
- In theory it should be a superset of loop-vectorizer
 - Unroll loop and vectorize with SLP
 - Even if loop-vectorizer fails, SLP could partly succeed

SLP Straight-Line Code Vectorizer

- Superword Level Parallelism [Larsen PLDI'00]
- State-of-the-art straight-line code vectorizer
- Implemented in most compilers (including GCC and LLVM)
- In theory it should be a superset of loop-vectorizer
 - Unroll loop and vectorize with SLP
 - Even if loop-vectorizer fails, SLP could partly succeed
- In practice it is missing features present in the Loop vectorizer (Interleaved Loads, Predication)

SLP Vectorization Algorithm

- Input is scalar IR

Scalar Code

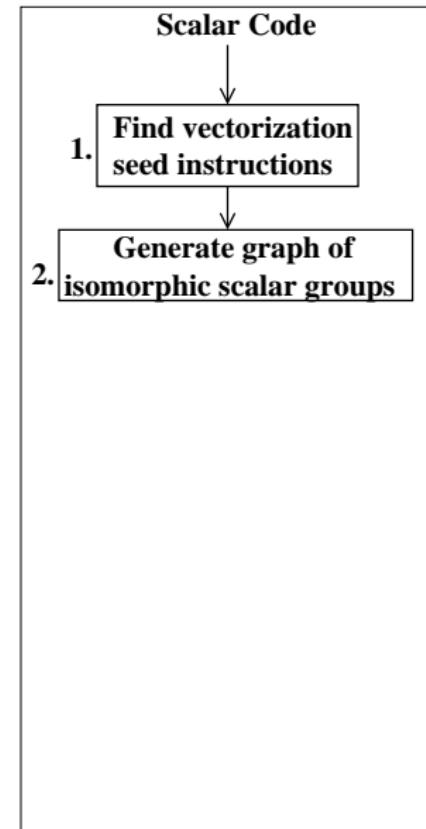
SLP Vectorization Algorithm

- Input is scalar IR
- Seed instructions are:
 - ① Consecutive Stores
 - ② Reductions



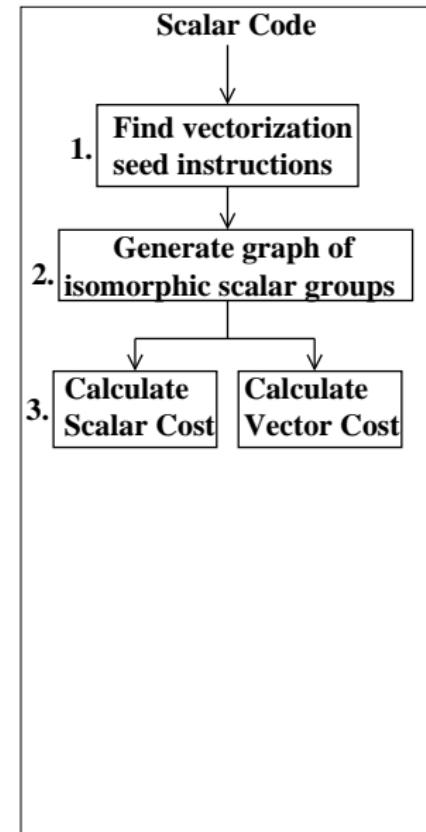
SLP Vectorization Algorithm

- Input is scalar IR
- Seed instructions are:
 - ① Consecutive Stores
 - ② Reductions
- Graph contains vectorizable isomorphic instructions



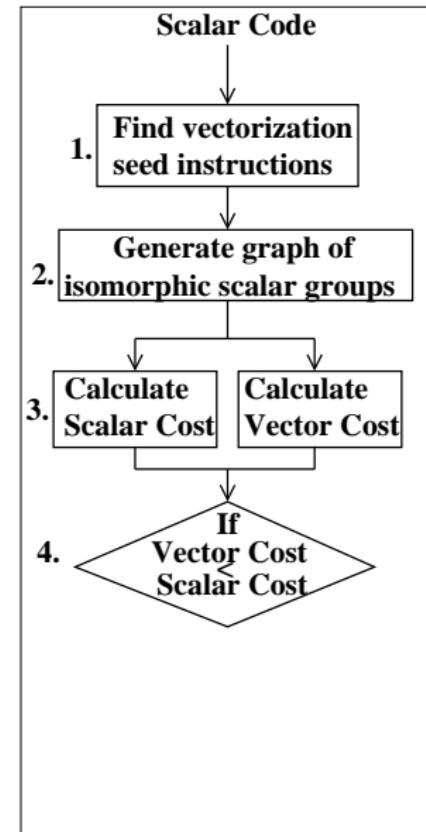
SLP Vectorization Algorithm

- Input is scalar IR
- Seed instructions are:
 - ① Consecutive Stores
 - ② Reductions
- Graph contains vectorizable isomorphic instructions
- Cost: weighted instr. count



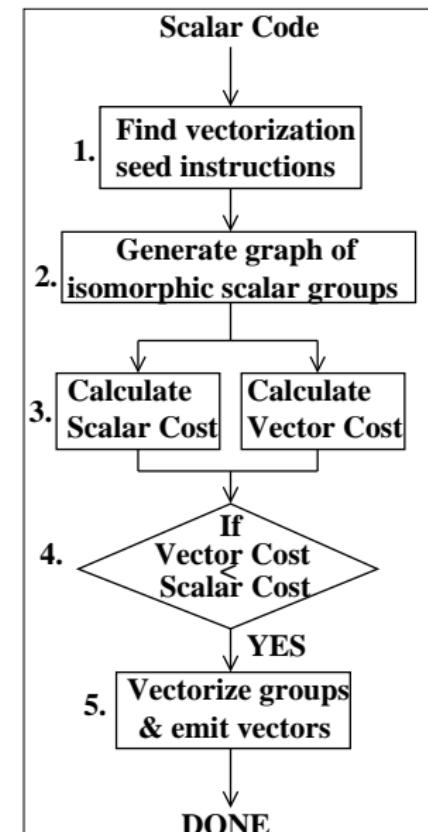
SLP Vectorization Algorithm

- Input is scalar IR
- Seed instructions are:
 - ① Consecutive Stores
 - ② Reductions
- Graph contains vectorizable isomorphic instructions
- Cost: weighted instr. count
- Check vectorization profitability



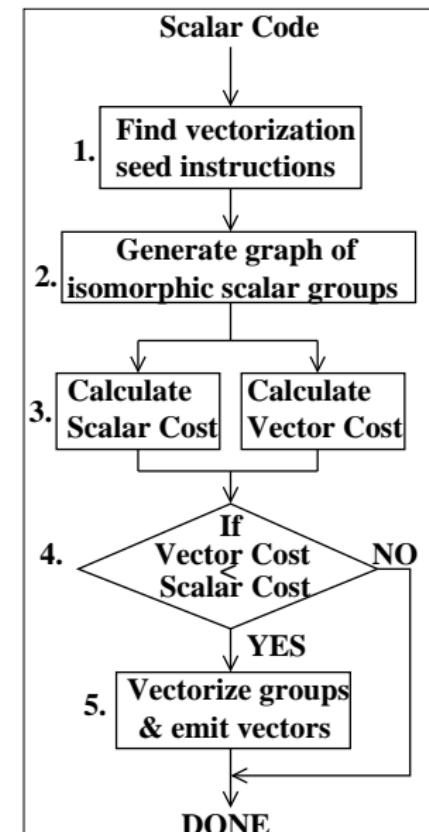
SLP Vectorization Algorithm

- Input is scalar IR
- Seed instructions are:
 - ① Consecutive Stores
 - ② Reductions
- Graph contains vectorizable isomorphic instructions
- Cost: weighted instr. count
- Check vectorization profitability
- Emit vectors only if profitable



SLP Vectorization Algorithm

- Input is scalar IR
- Seed instructions are:
 - ① Consecutive Stores
 - ② Reductions
- Graph contains vectorizable isomorphic instructions
- Cost: weighted instr. count
- Check vectorization profitability
- Emit vectors only if profitable



SLP Fails due to non-isomorphism

```
...
B[i] = A[i] * 7.0 + 1.0;
B[i+1]= A[i+1]      + 5.0;
...
```

a. Input C code

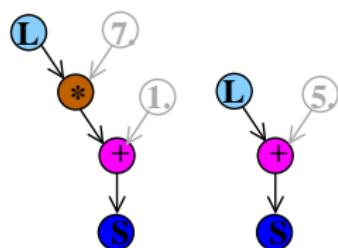
(X) Instruction Node or Constant → Data Flow Edge

SLP Fails due to non-isomorphism

```

    . . .
B[i] = A[i] * 7.0 + 1.0;
B[i+1] = A[i+1] + 5.0;
    . . .
    
```

a. Input C code



b. DFG

	Instruction Node or Constant	→ Data Flow Edge
--	------------------------------	------------------

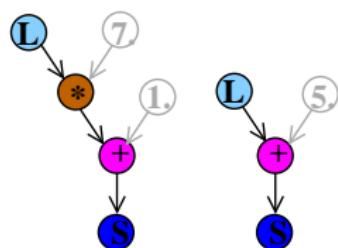
SLP Fails due to non-isomorphism

```

B[i] = A[i] * 7.0 + 1.0;
B[i+1] = A[i+1] + 5.0;
...

```

a. Input C code



b. DFG



c. SLP internal graph



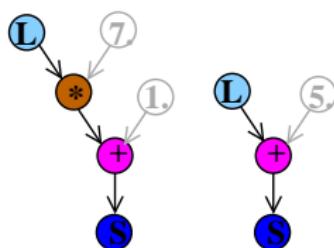
d. SLP vectorized groups

	Instruction Node or Constant	→ Data Flow Edge
--	------------------------------	------------------

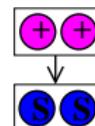
SLP Fails due to non-isomorphism

$$\begin{array}{l} \dots \\ B[i] = A[i] * 7.0 + 1.0; \\ B[i+1] = A[i+1] + 5.0; \\ \dots \end{array}$$

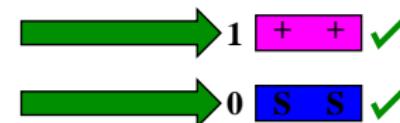
a. Input C code



b. DFG



c. SLP internal graph



d. SLP vectorized groups

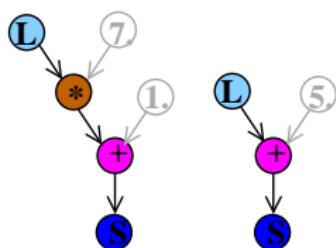
(X) Instruction Node or Constant → Data Flow Edge

SLP Fails due to non-isomorphism

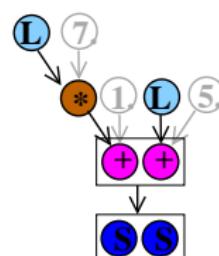
```

    . . .
B[i] = A[i] * 7.0 + 1.0;
B[i+1] = A[i+1] + 5.0;
    . . .
    
```

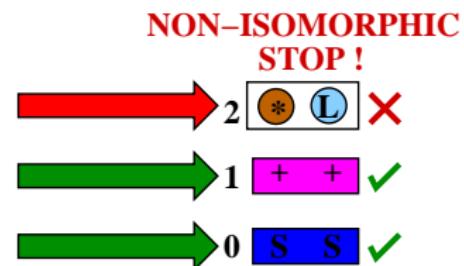
a. Input C code



b. DFG



c. SLP internal graph



d. SLP vectorized groups

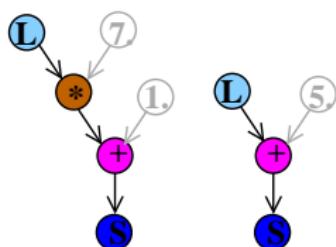
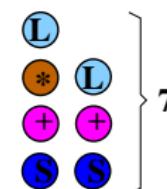
(X) Instruction Node or Constant → Data Flow Edge

SLP Fails due to non-isomorphism

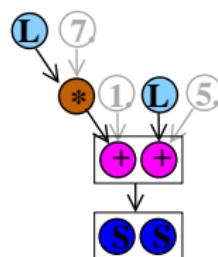
\dots
 $B[i] = A[i] * 7.0 + 1.0;$
 $B[i+1] = A[i+1] + 5.0;$
 \dots

a. Input C code

Scalar Cost

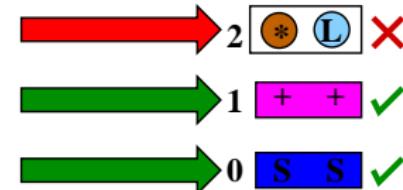


b. DFG



c. SLP internal graph

**NON-ISOMORPHIC
STOP!**



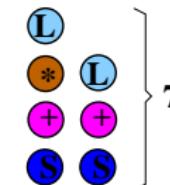
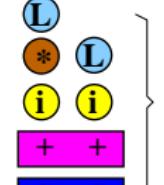
d. SLP vectorized groups

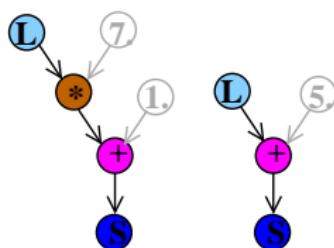
(X) Instruction Node or Constant → Data Flow Edge

SLP Fails due to non-isomorphism

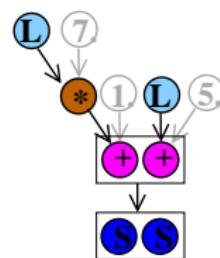
$B[i] = A[i] * 7.0 + 1.0;$
 $B[i+1] = A[i+1] + 5.0;$
 \dots

a. Input C code

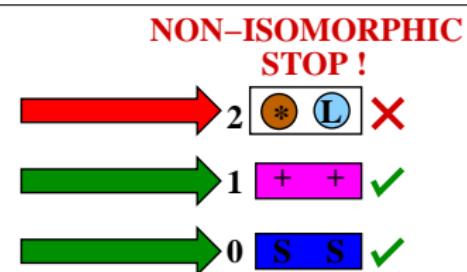
Scalar Cost	Vector Cost
	



b. DFG



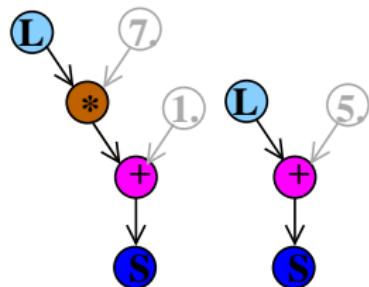
c. SLP internal graph



d. SLP vectorized groups

X Instruction Node or Constant → Data Flow Edge

PSLP fixes Non-Isomorphism



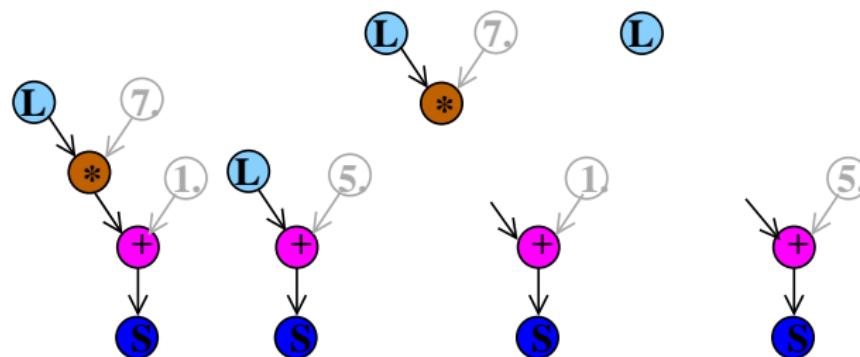
a. PSLP graphs



Instruction or Constant

→Data Flow Edge

PSLP fixes Non-Isomorphism



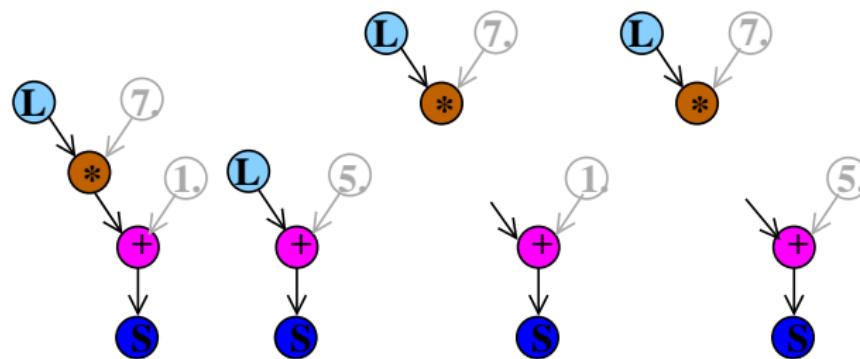
a. PSLP graphs

b. PSLP padded graphs

(X) Instruction or Constant

→Data Flow Edge

PSLP fixes Non-Isomorphism



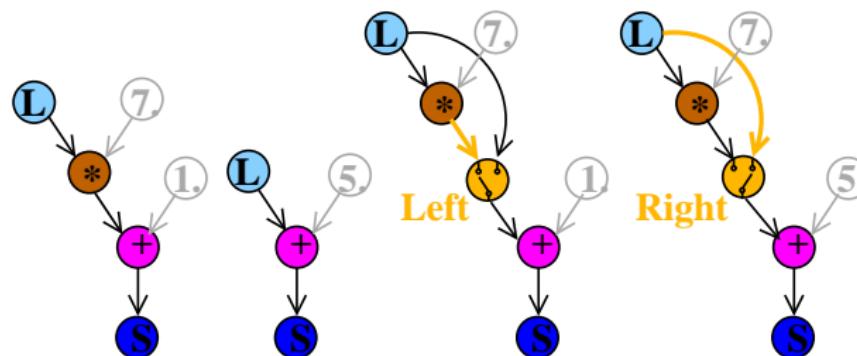
a. PSLP graphs

b. PSLP padded graphs

(X) Instruction or Constant

→Data Flow Edge

PSLP fixes Non-Isomorphism

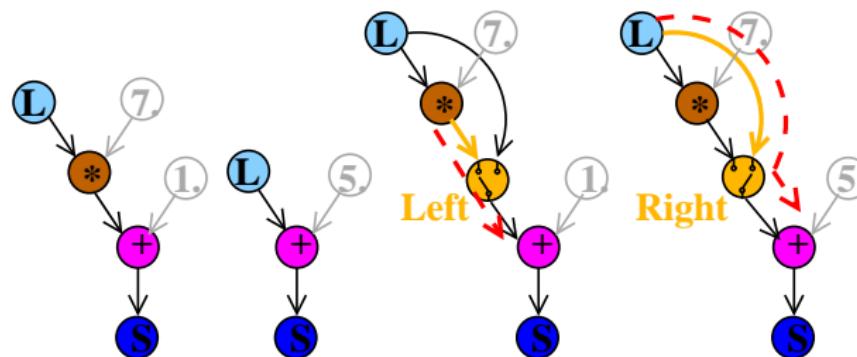


a. PSLP graphs

b. PSLP padded graphs

 Instruction or Constant	 Select Instruction	→Data Flow Edge
--	--	-----------------

PSLP fixes Non-Isomorphism

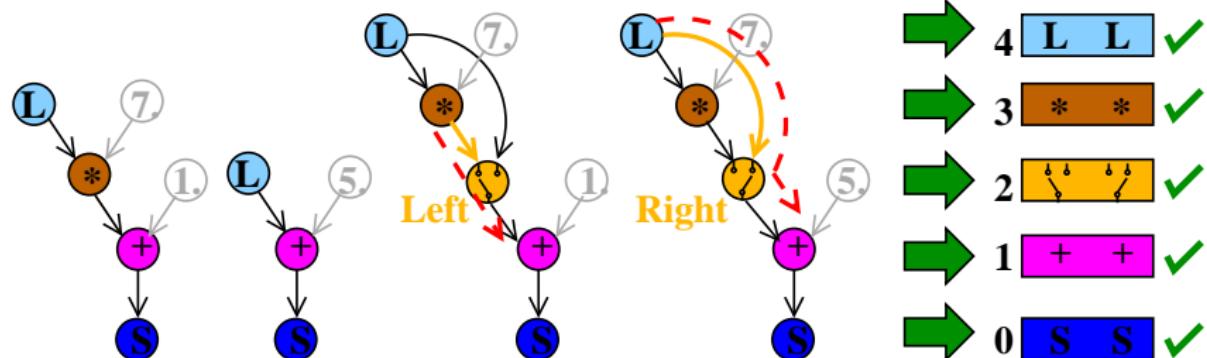


a. PSLP graphs

b. PSLP padded graphs

 Instruction or Constant	 Select Instruction	→ Data Flow Edge
--	--	------------------

PSLP fixes Non-Isomorphism



a. PSLP graphs

b. PSLP padded graphs

c. PSLP groups



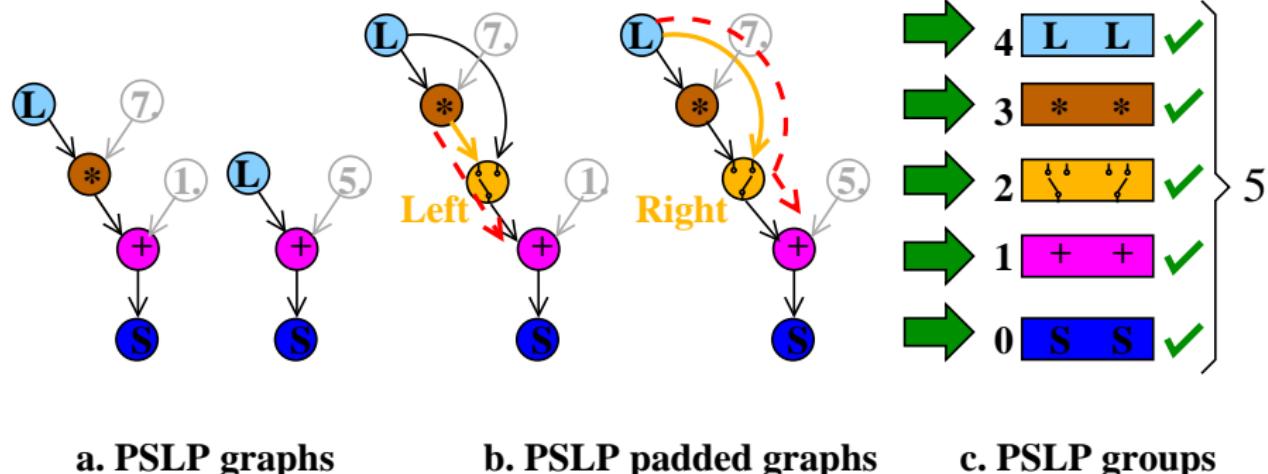
Instruction or Constant



Select Instruction

→ Data Flow Edge

PSLP fixes Non-Isomorphism



a. PSLP graphs

b. PSLP padded graphs

c. PSLP groups



Instruction or Constant



Select Instruction

→ Data Flow Edge

PSLP Algorithm

- Extension to SLP

PSLP Algorithm

- Extension to SLP

1. Find vectorization seed instructions

PSLP Algorithm

- Extension to SLP
- Generate multiple graphs (unlike SLP)

1. Find vectorization seed instructions
2. Generate a graph for each seed

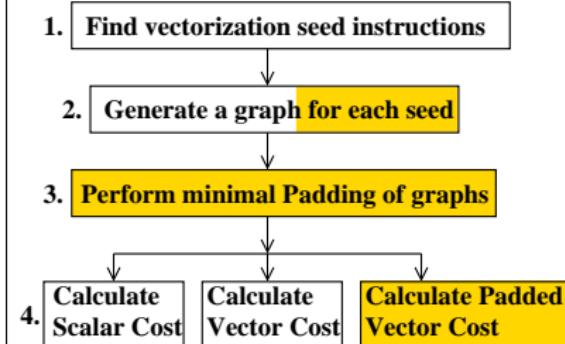
PSLP Algorithm

- Extension to SLP
- Generate multiple graphs (unlike SLP)
- Minimal Padding

1. Find vectorization seed instructions
2. Generate a graph for each seed
3. Perform minimal Padding of graphs

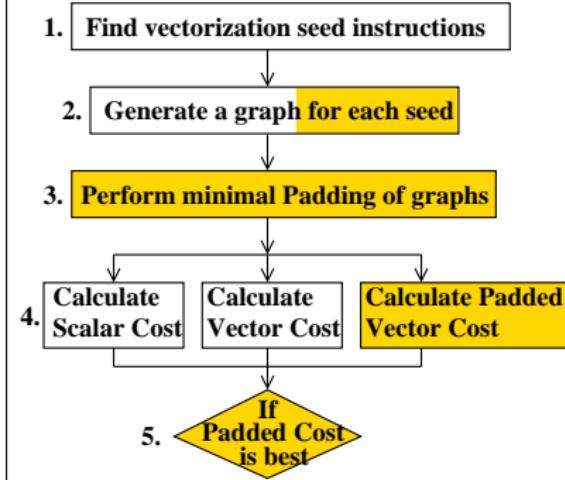
PSLP Algorithm

- Extension to SLP
- Generate multiple graphs (unlike SLP)
- Minimal Padding
- Cost estimation



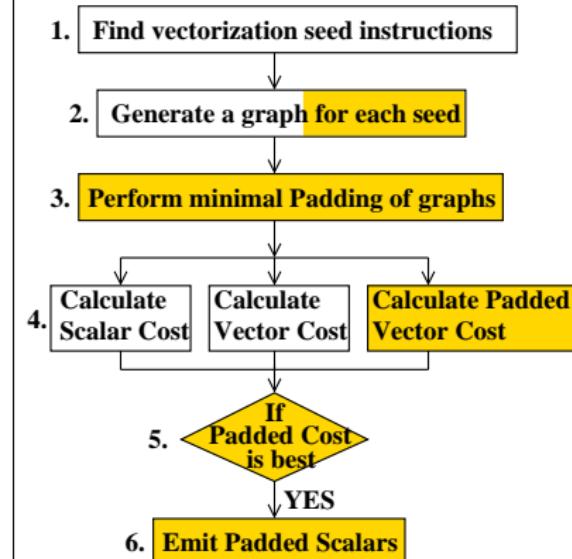
PSLP Algorithm

- Extension to SLP
- Generate multiple graphs (unlike SLP)
- Minimal Padding
- Cost estimation



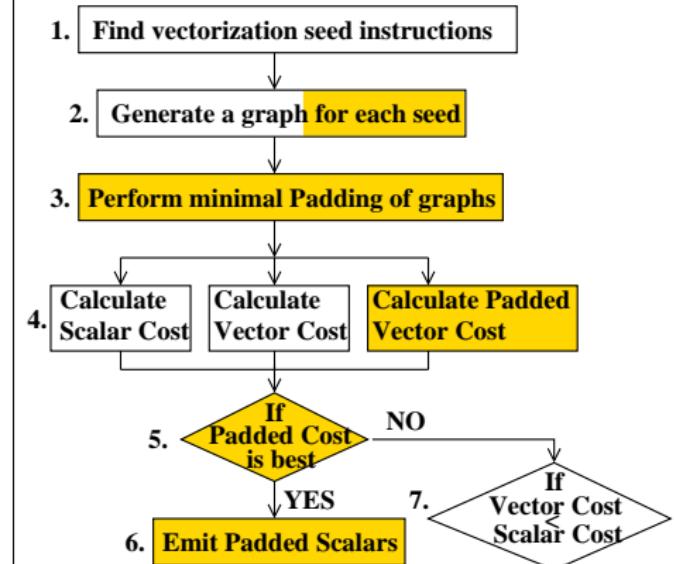
PSLP Algorithm

- Extension to SLP
- Generate multiple graphs (unlike SLP)
- Minimal Padding
- Cost estimation
- Emit redundant code to create isomorphism



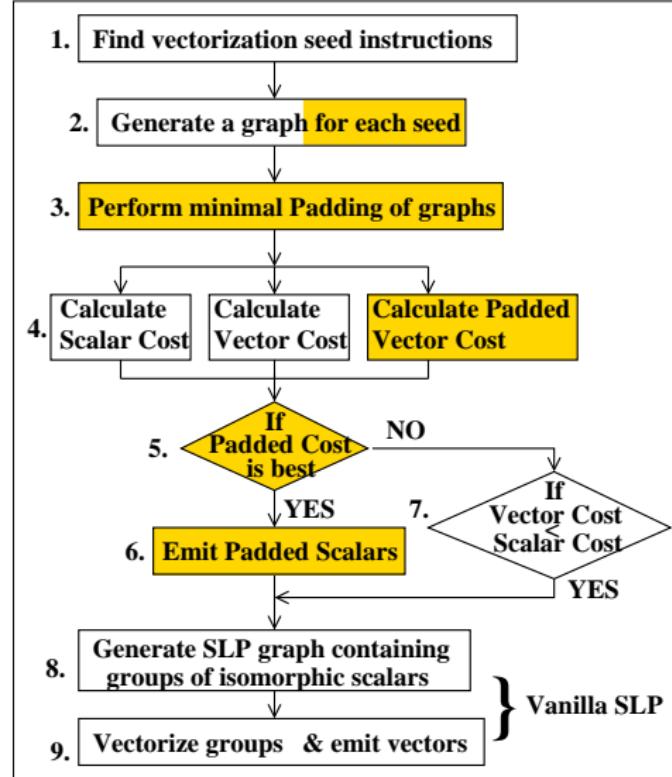
PSLP Algorithm

- Extension to SLP
- Generate multiple graphs (unlike SLP)
- Minimal Padding
- Cost estimation
- Emit redundant code to create isomorphism



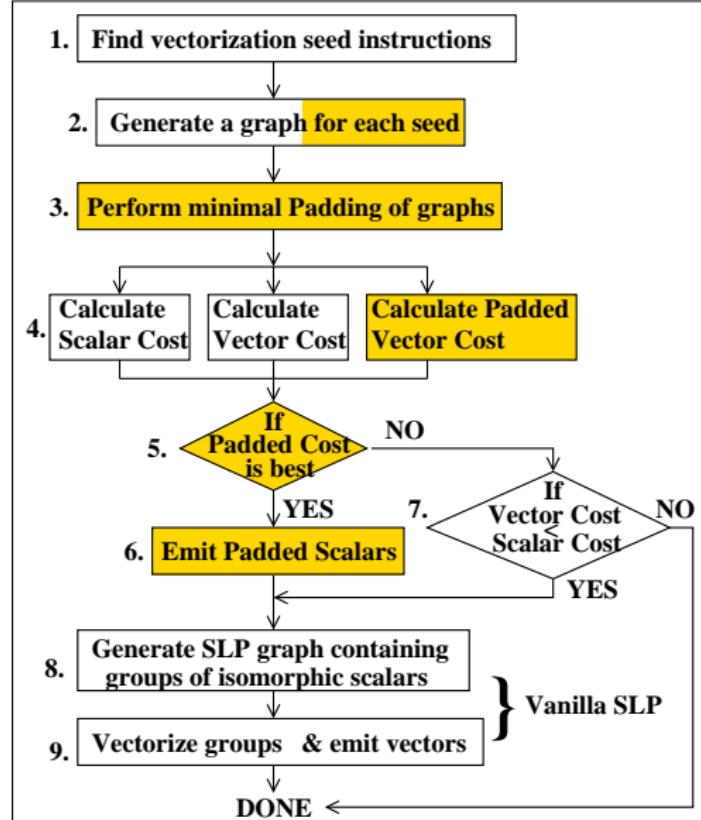
PSLP Algorithm

- Extension to SLP
- Generate multiple graphs (unlike SLP)
- Minimal Padding
- Cost estimation
- Emit redundant code to create isomorphism
- Code vectorized by original SLP

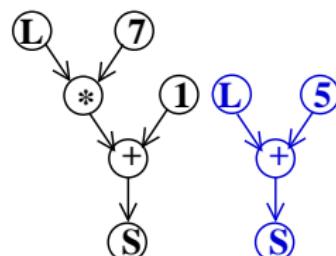
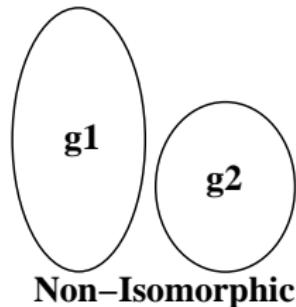


PSLP Algorithm

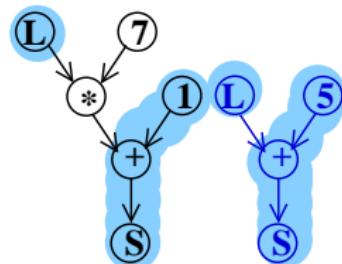
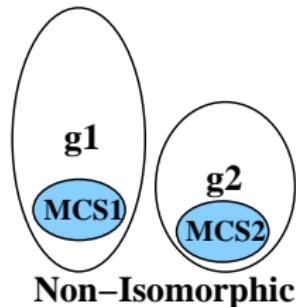
- Extension to SLP
 - Generate multiple graphs (unlike SLP)
 - Minimal Padding
 - Cost estimation
 - Emit redundant code to create isomorphism
 - Code vectorized by original SLP



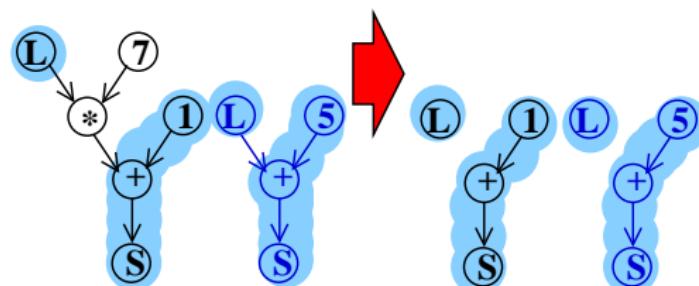
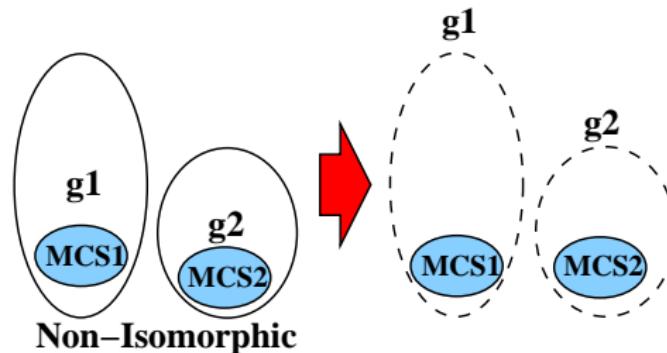
Minimal Padding Algorithm



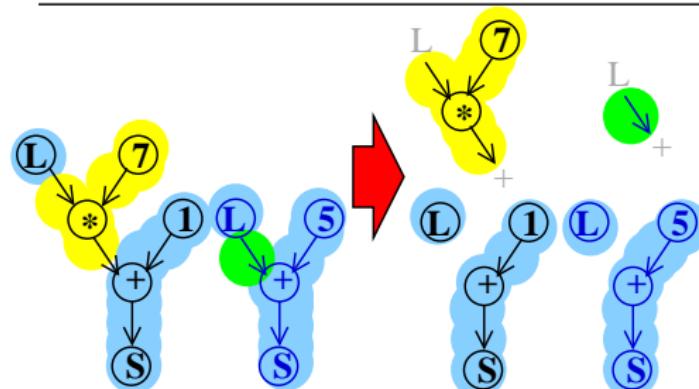
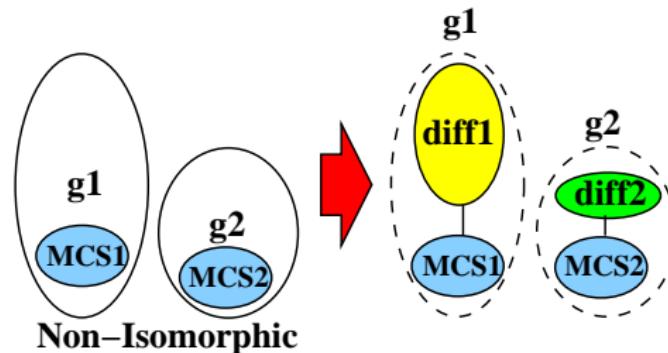
Minimal Padding Algorithm



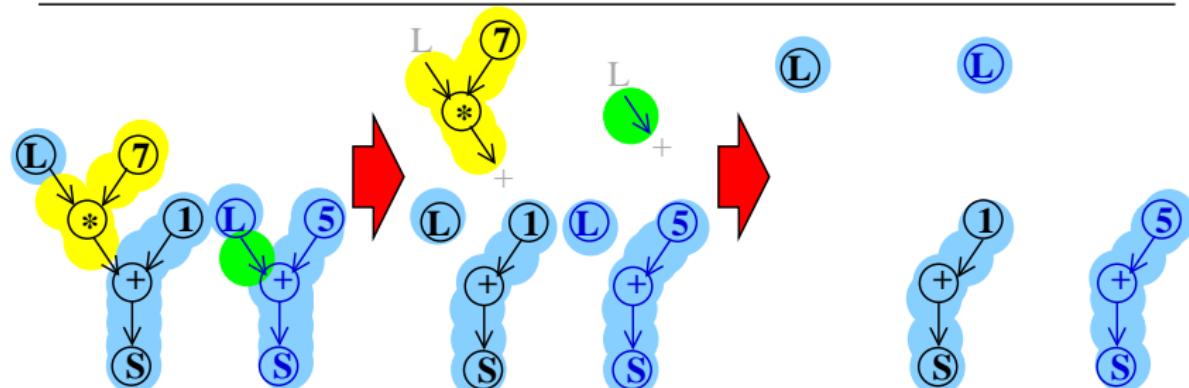
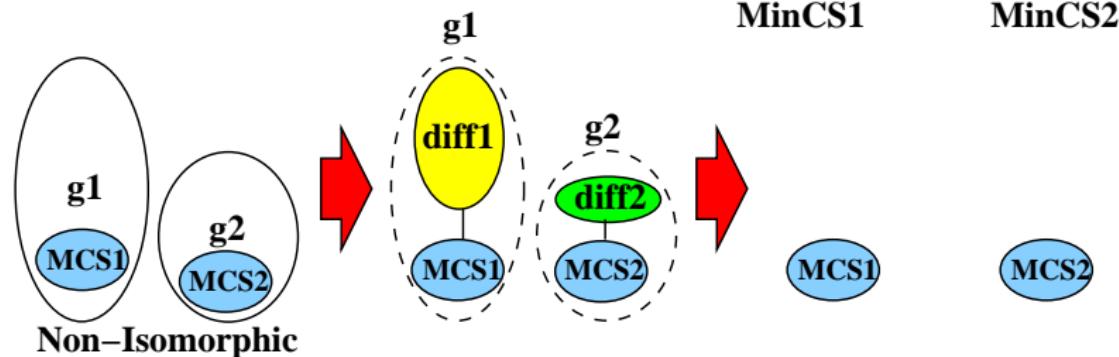
Minimal Padding Algorithm



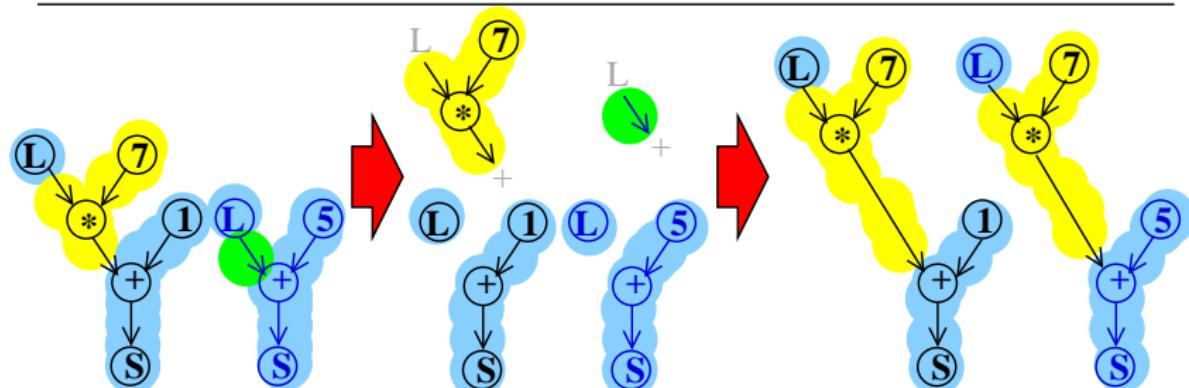
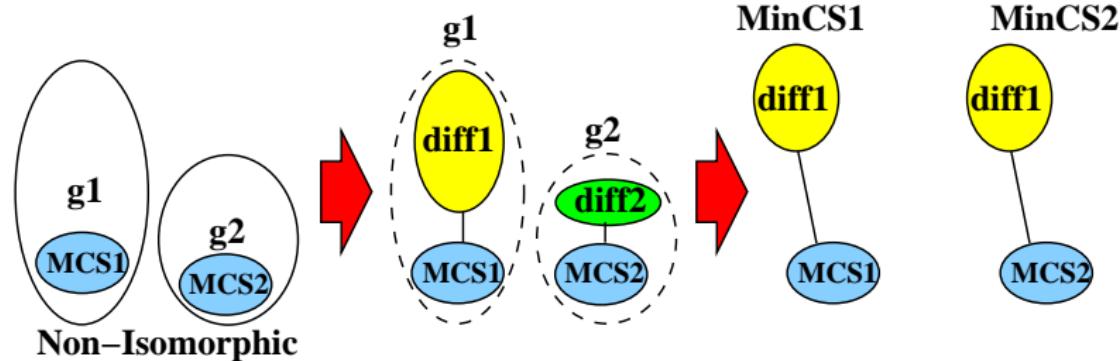
Minimal Padding Algorithm



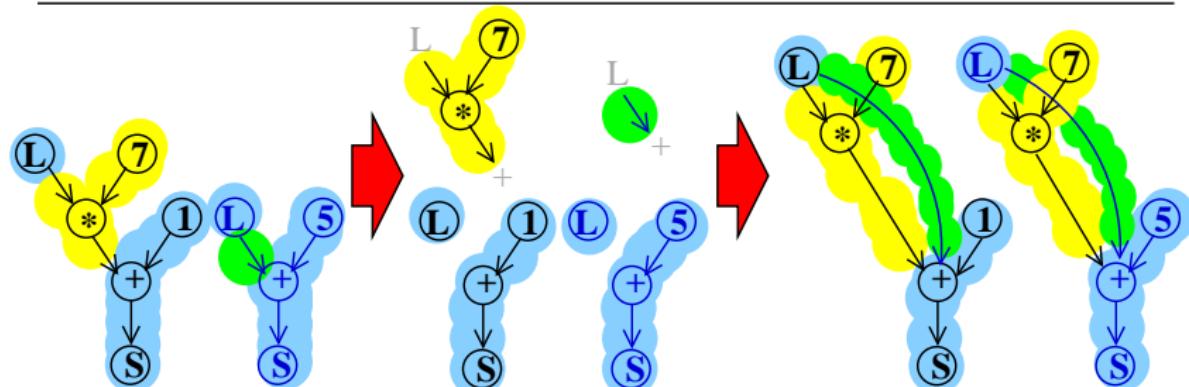
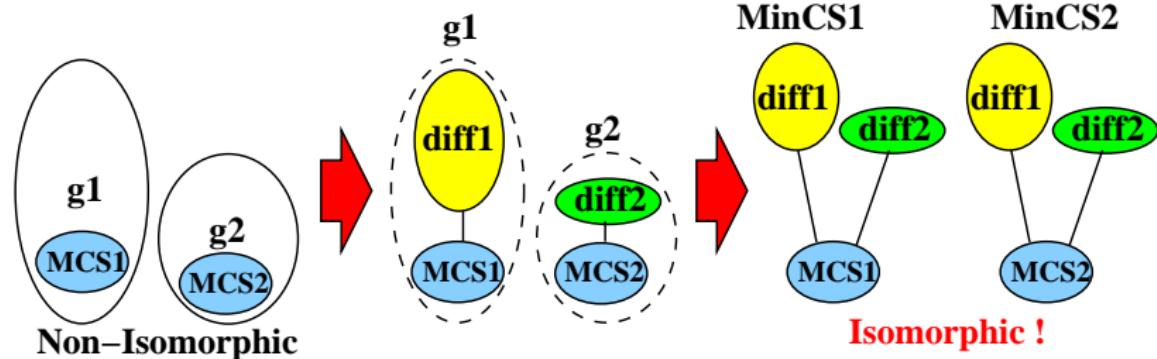
Minimal Padding Algorithm



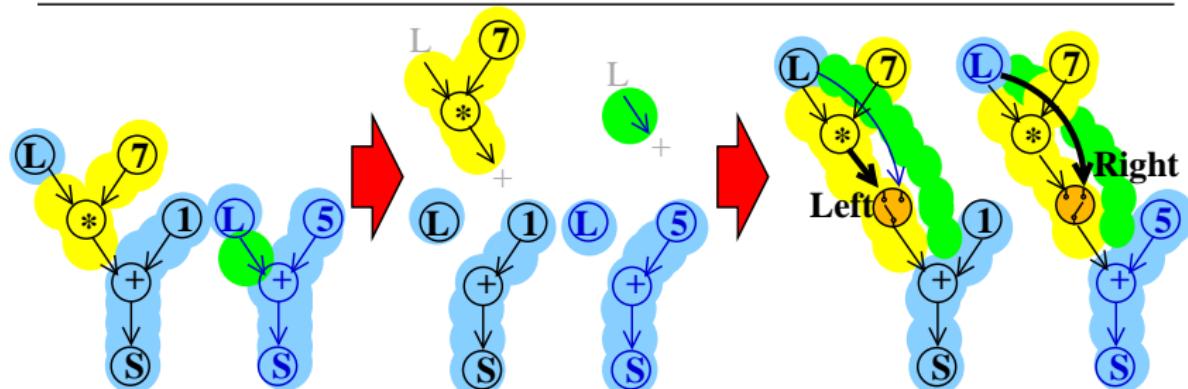
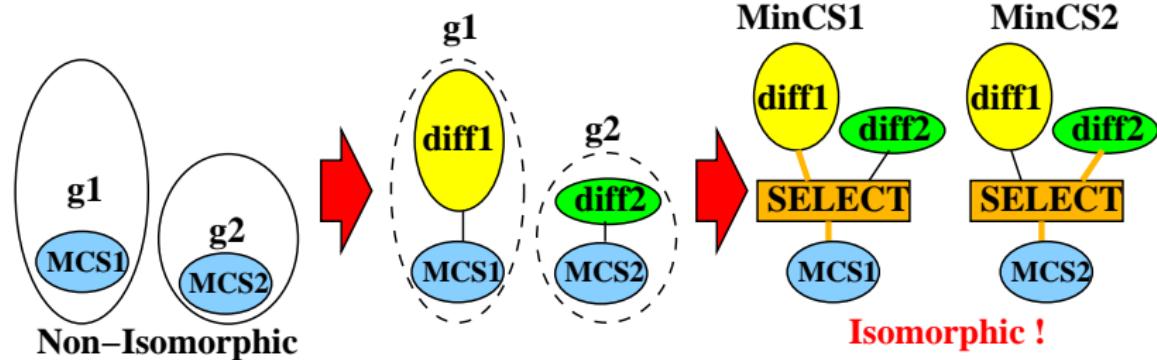
Minimal Padding Algorithm



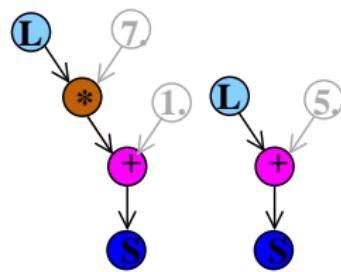
Minimal Padding Algorithm



Minimal Padding Algorithm

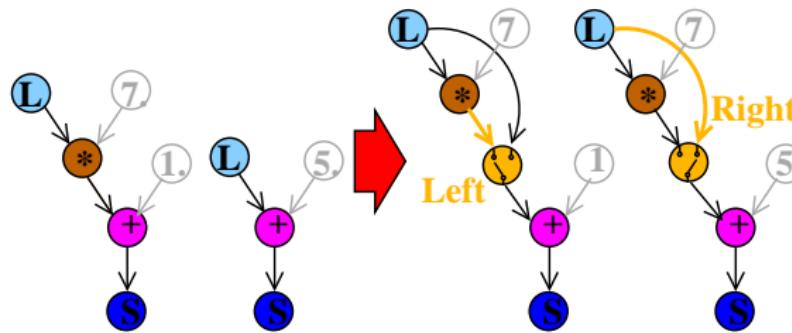


We can do better: Remove redundant Selects



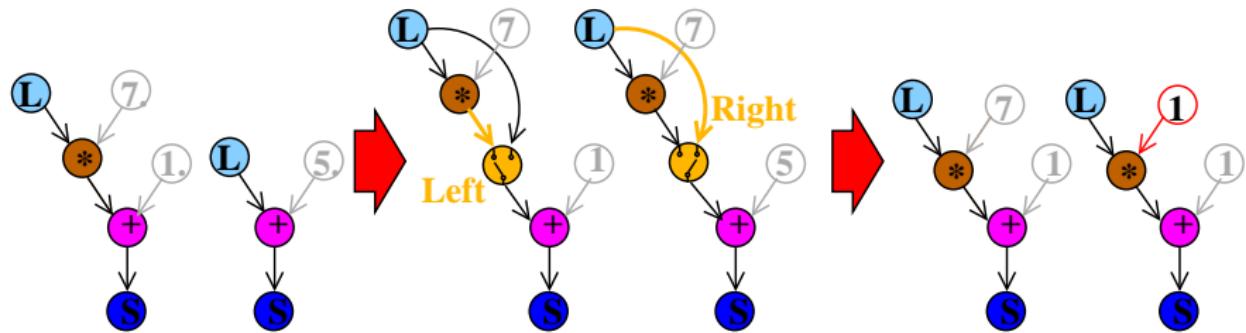
EXAMPLE: Instruction acting as Select

We can do better: Remove redundant Selects



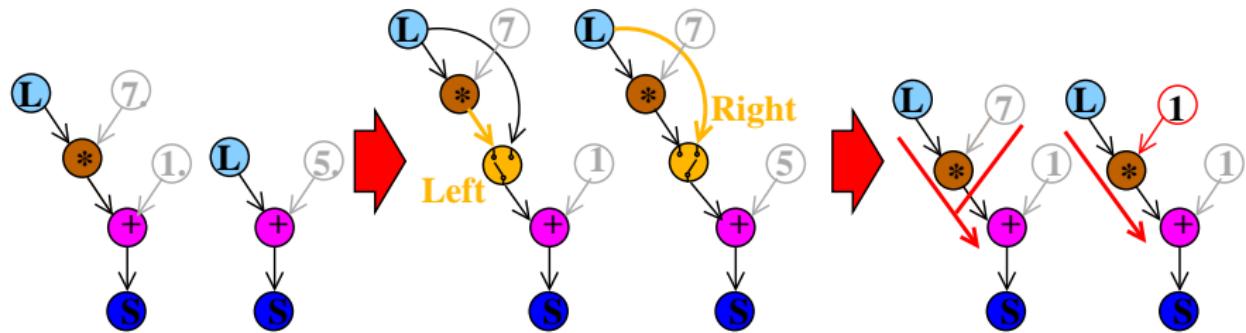
EXAMPLE: Instruction acting as Select

We can do better: Remove redundant Selects



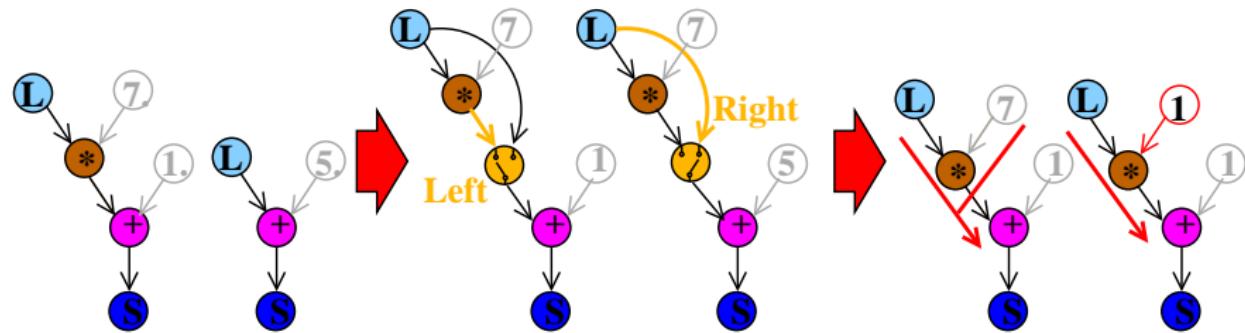
EXAMPLE: Instruction acting as Select

We can do better: Remove redundant Selects



EXAMPLE: Instruction acting as Select

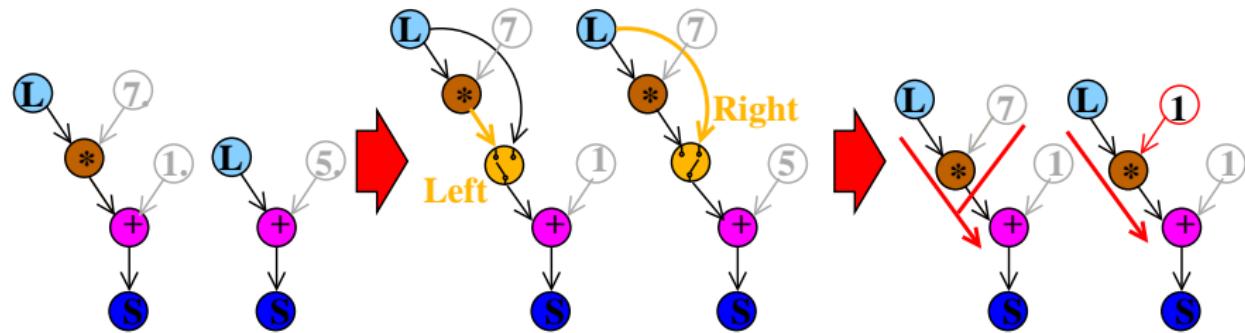
We can do better: Remove redundant Selects



EXAMPLE: Instruction acting as Select



We can do better: Remove redundant Selects

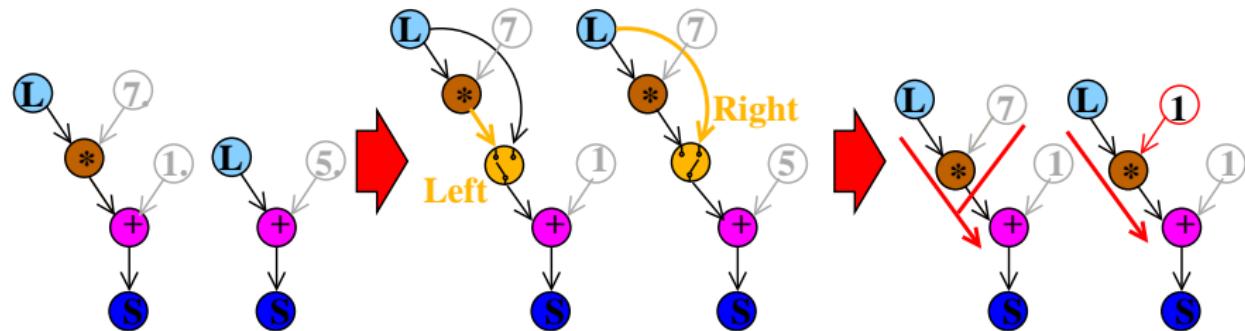


EXAMPLE: Instruction acting as Select

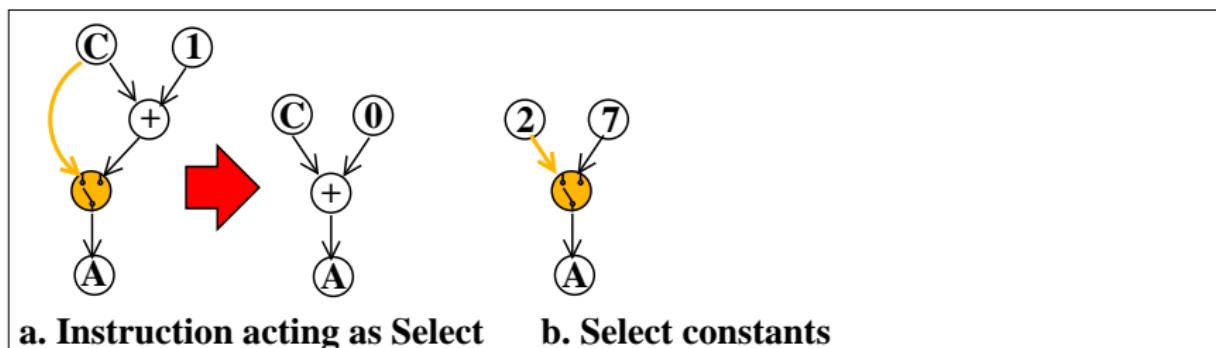


a. Instruction acting as Select

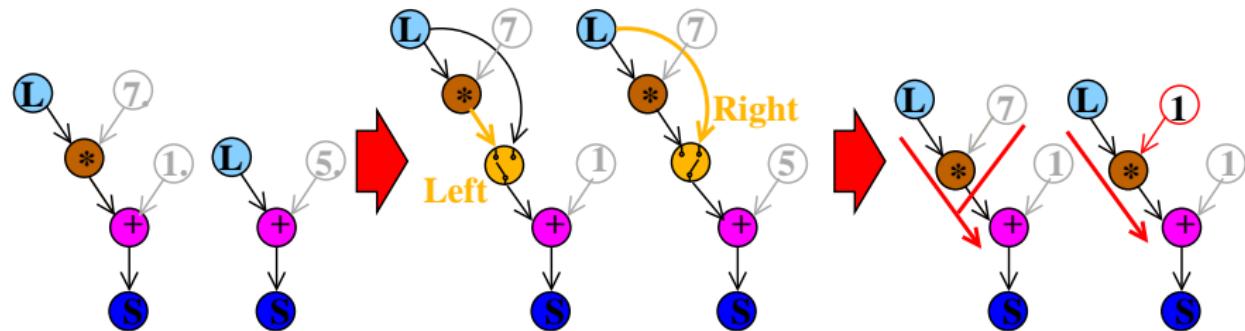
We can do better: Remove redundant Selects



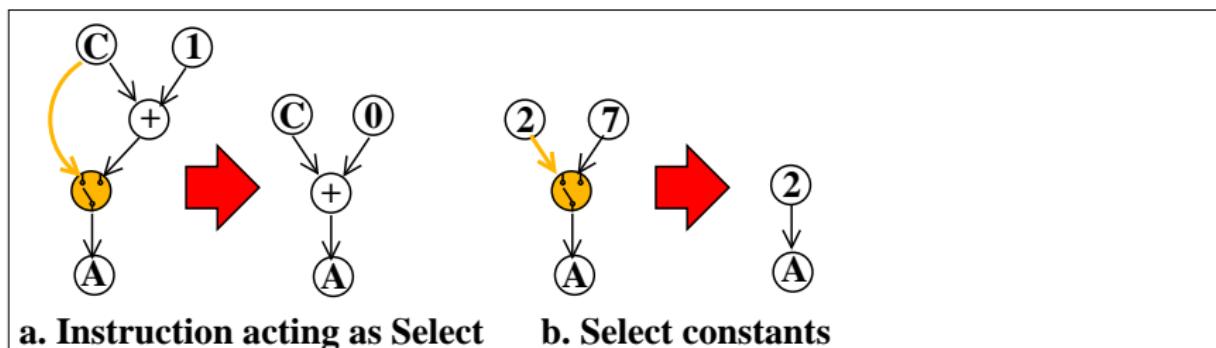
EXAMPLE: Instruction acting as Select



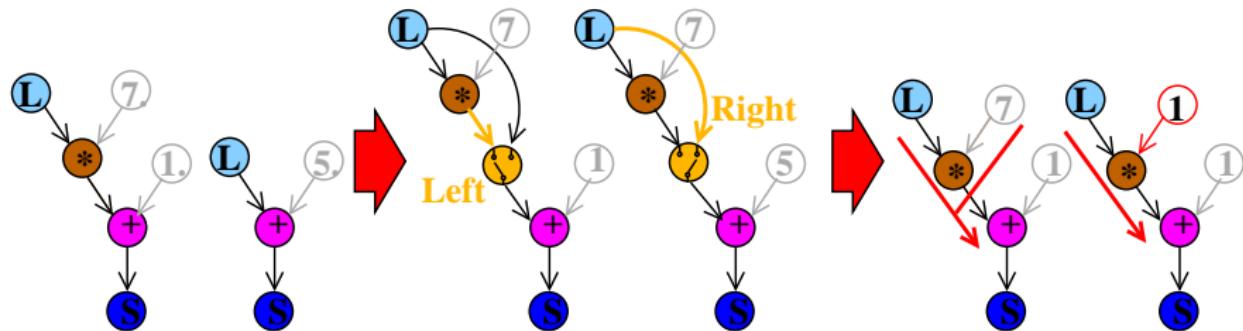
We can do better: Remove redundant Selects



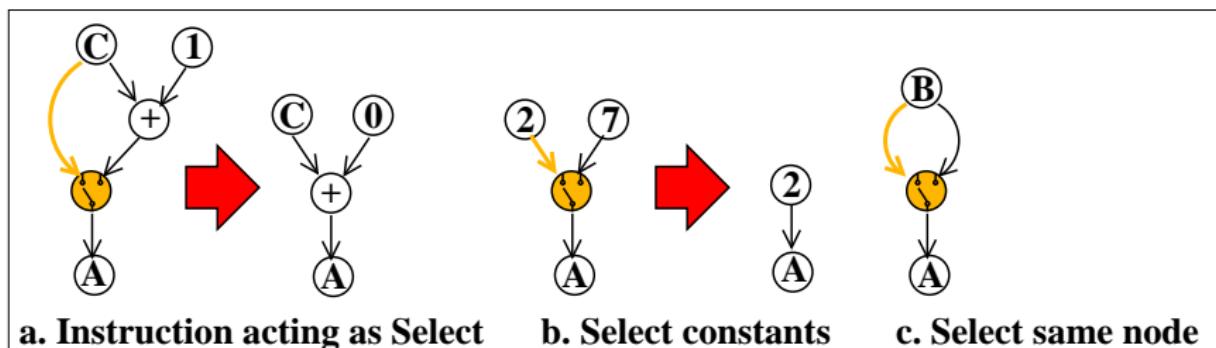
EXAMPLE: Instruction acting as Select



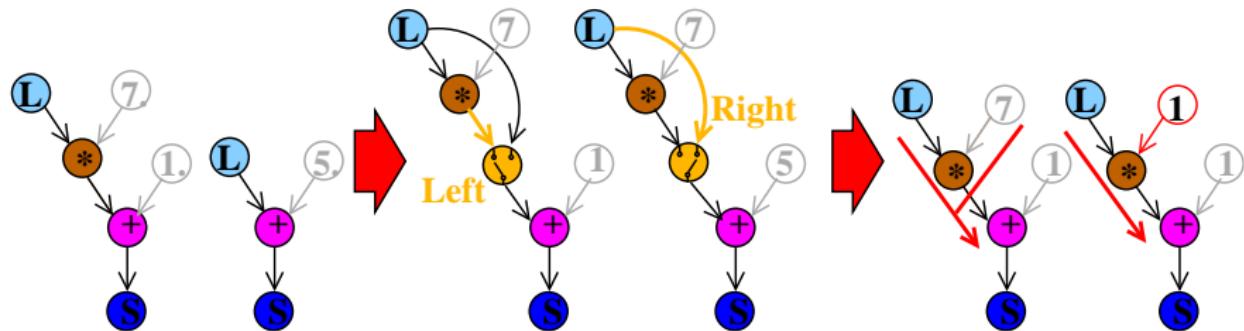
We can do better: Remove redundant Selects



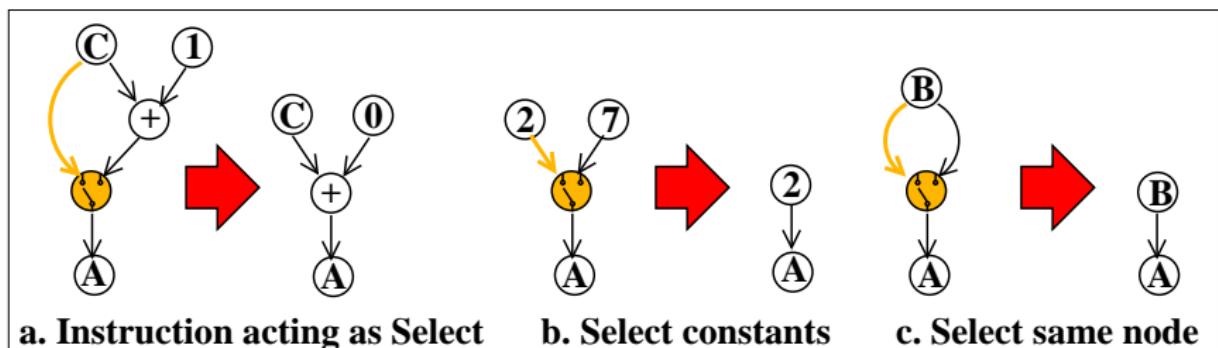
EXAMPLE: Instruction acting as Select



We can do better: Remove redundant Selects



EXAMPLE: Instruction acting as Select



a. Instruction acting as Select

b. Select constants

c. Select same node

Opportunities for PSLP in real-life applications

- ① Non-isomorphic source code (e.g. computing conjugates in 433.milc)

```
b[0].real = a[0].real
b[0].imag = - a[0].imag
b[1].real = a[1].real
b[1].imag = - a[1].imag
```

...	Memory
a[0].real	
a[0].imag	
a[1].real	
a[1].imag	
...	

Opportunities for PSLP in real-life applications

- ① Non-isomorphic source code (e.g. computing conjugates in 433.milc)

```
b[0].real = a[0].real
b[0].imag =  a[0].imag
b[1].real = a[1].real
b[1].imag =  a[1].imag
```

...	Memory
a[0].real	
a[0].imag	
a[1].real	
a[1].imag	
...	

Opportunities for PSLP in real-life applications

- ① Non-isomorphic source code (e.g. computing conjugates in 433.milc)

```
b[0].real = a[0].real
b[0].imag =  a[0].imag
b[1].real = a[1].real
b[1].imag =  a[1].imag
```

Memory
...
a[0].real
a[0].imag
a[1].real
a[1].imag
...

- ② Isomorphic source code but non-isomorphic IR due to high-level optimizations (jdct of jpeg)

```
tmp1 = quantval[0]*16384
tmp2 = quantval[1]*22725
tmp3 = quantval[2]*21407
tmp4 = quantval[3]*19266
```

Opportunities for PSLP in real-life applications

- ① Non-isomorphic source code (e.g. computing conjugates in 433.milc)

```
b[0].real = a[0].real
b[0].imag =  a[0].imag
b[1].real = a[1].real
b[1].imag =  a[1].imag
```

...	Memory
a[0].real	
a[0].imag	
a[1].real	
a[1].imag	
...	

- ② Isomorphic source code but non-isomorphic IR due to high-level optimizations (jdct of jpeg)

```
tmp1 = quantval[0]*16384
tmp2 = quantval[1]*22725
tmp3 = quantval[2]*21407
tmp4 = quantval[3]*19266
```



```
tmp1 = quantval[0]<<14
tmp2 = quantval[1]*22725
tmp3 = quantval[2]*21407
tmp4 = quantval[3]*19266
```

Opportunities for PSLP in real-life applications

- ① Non-isomorphic source code (e.g. computing conjugates in 433.milc)

```
b[0].real = a[0].real
b[0].imag =  a[0].imag
b[1].real = a[1].real
b[1].imag =  a[1].imag
```

...	Memory
a[0].real	
a[0].imag	
a[1].real	
a[1].imag	
...	

- ② Isomorphic source code but non-isomorphic IR due to high-level optimizations (jdct of jpeg)

```
tmp1 = quantval[0]*16384
tmp2 = quantval[1]*22725
tmp3 = quantval[2]*21407
tmp4 = quantval[3]*19266
```



```
tmp1 = quantval[0]<<14
tmp2 = quantval[1]*22725
tmp3 = quantval[2]*21407
tmp4 = quantval[3]*19266
```

Experimental Setup

- Implemented PSLP in the trunk version of the LLVM 3.6 compiler.

Experimental Setup

- Implemented PSLP in the trunk version of the LLVM 3.6 compiler.
- Target: Intel Core i5-4570 @ 3.2Ghz

Experimental Setup

- Implemented PSLP in the trunk version of the LLVM 3.6 compiler.
- Target: Intel Core i5-4570 @ 3.2Ghz
- Compiler flags: -O3 -allow-partial-unroll
-march=core-avx2 -mtune-core-i7 -ffast-math

Experimental Setup

- Implemented PSLP in the trunk version of the LLVM 3.6 compiler.
- Target: Intel Core i5-4570 @ 3.2Ghz
- Compiler flags: -O3 -allow-partial-unroll
-march=core-avx2 -mtune-core-i7 -ffast-math
- Kernels, SPEC 2006 and Mediabench II
- We evaluated the following cases:

Experimental Setup

- Implemented PSLP in the trunk version of the LLVM 3.6 compiler.
- Target: Intel Core i5-4570 @ 3.2Ghz
- Compiler flags: -O3 -allow-partial-unroll
-march=core-avx2 -mtune-core-i7 -ffast-math
- Kernels, SPEC 2006 and Mediabench II
- We evaluated the following cases:
 - ① All loop, SLP and PSLP vectorizers disabled (O3)

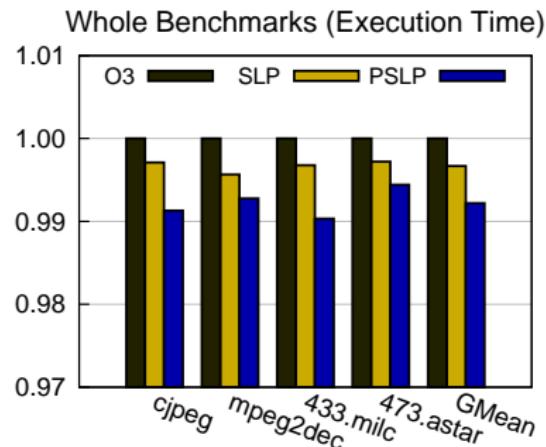
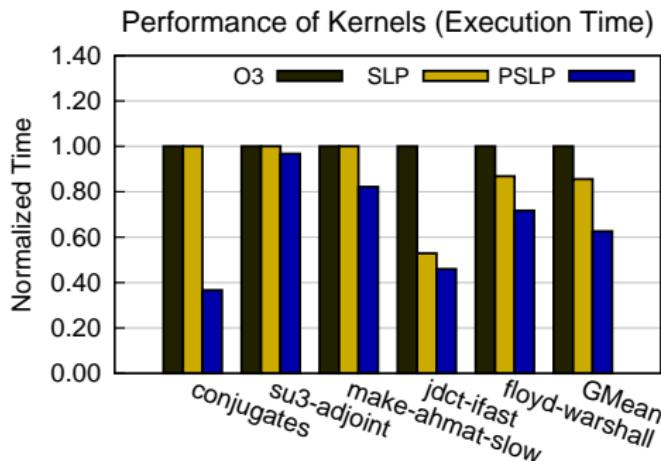
Experimental Setup

- Implemented PSLP in the trunk version of the LLVM 3.6 compiler.
- Target: Intel Core i5-4570 @ 3.2Ghz
- Compiler flags: -O3 -allow-partial-unroll
-march=core-avx2 -mtune-core-i7 -ffast-math
- Kernels, SPEC 2006 and Mediabench II
- We evaluated the following cases:
 - ① All loop, SLP and PSLP vectorizers disabled (O3)
 - ② O3 + SLP enabled (**SLP**)

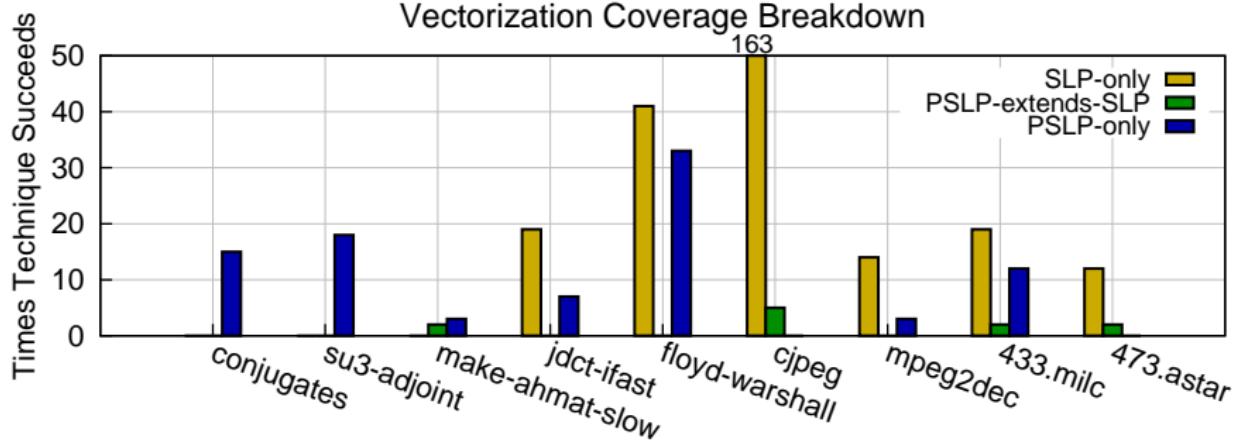
Experimental Setup

- Implemented PSLP in the trunk version of the LLVM 3.6 compiler.
- Target: Intel Core i5-4570 @ 3.2Ghz
- Compiler flags: -O3 -allow-partial-unroll
-march=core-avx2 -mtune-core-i7 -ffast-math
- Kernels, SPEC 2006 and Mediabench II
- We evaluated the following cases:
 - ① All loop, SLP and PSLP vectorizers disabled (O3)
 - ② O3 + SLP enabled (**SLP**)
 - ③ O3 + PSLP enabled (**PSLP**)

PSLP increases performance



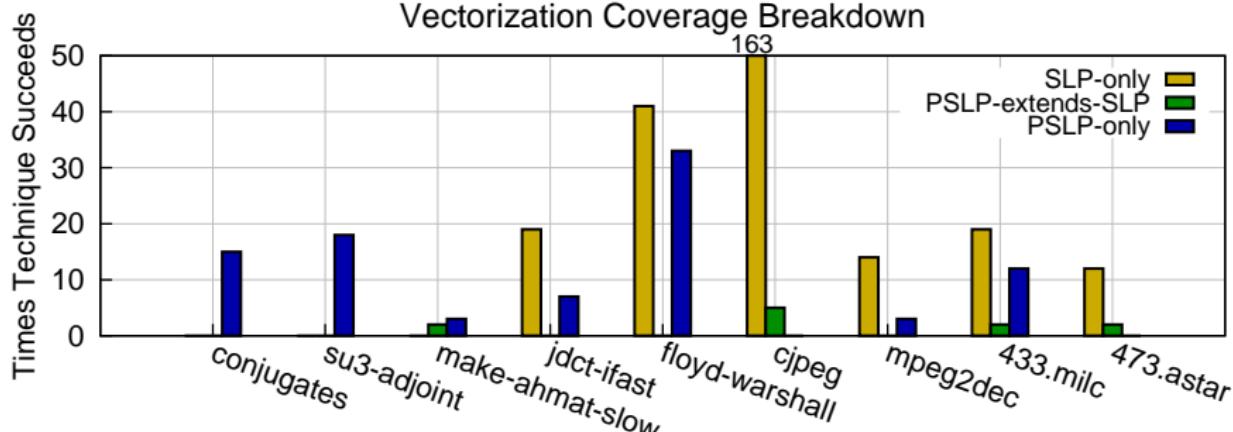
PSLP enables or extends vectorization



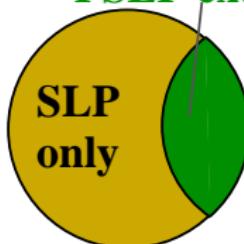
- SLP is adequate


 SLP
only

PSLP enables or extends vectorization

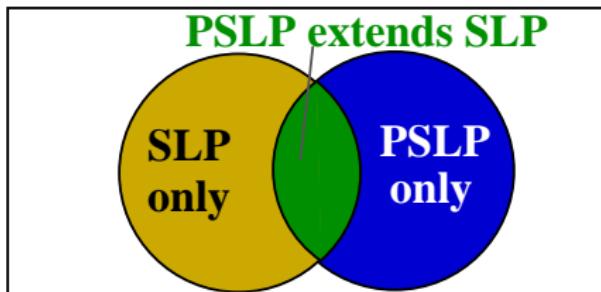
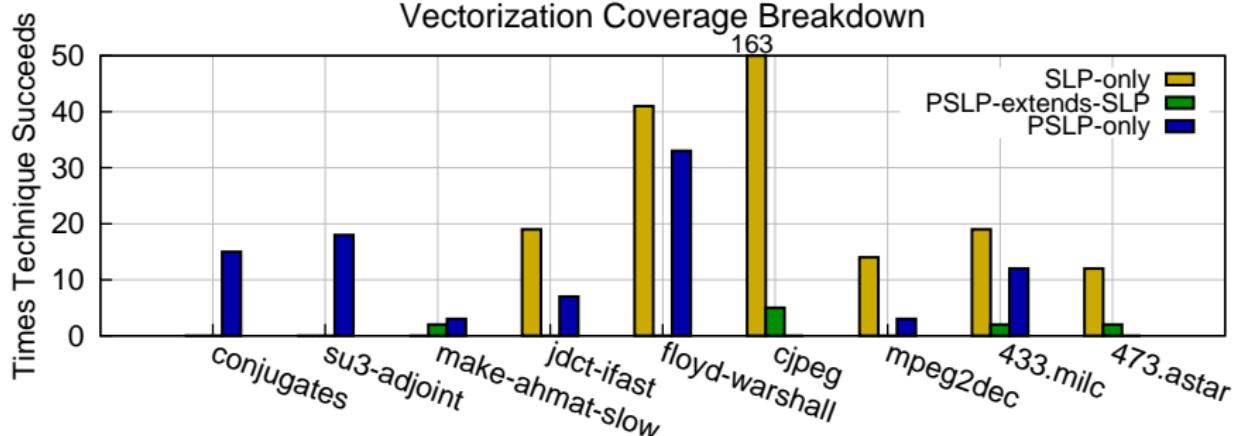


PSLP extends SLP



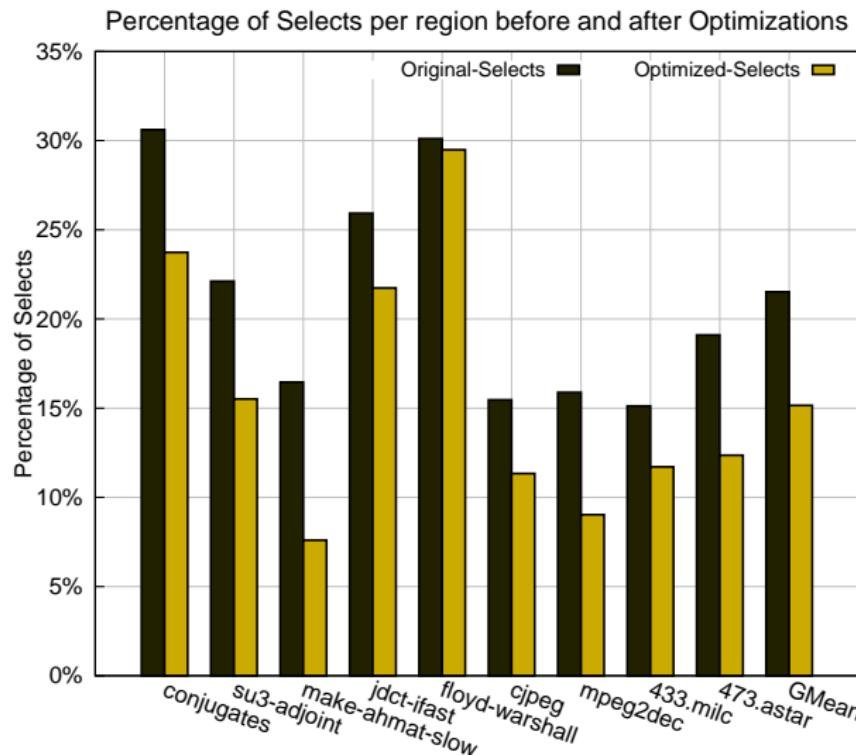
- SLP is adequate
- SLP stops at non-isomorphic code. PSLP extends it.

PSLP enables or extends vectorization



- SLP is adequate
- SLP stops at non-isomorphic code. PSLP extends it.
- SLP fails completely. PSLP succeeds.

Optimizing away redundant Selects



- Select-removal optimizations remove about 21% of the Selects

Conclusion

- PSLP improves vectorization coverage compared to the state-of-the-art

Conclusion

- PSLP improves vectorization coverage compared to the state-of-the-art
- Converts non-isomorphic code into isomorphic by:

Conclusion

- PSLP improves vectorization coverage compared to the state-of-the-art
- Converts non-isomorphic code into isomorphic by:
 - Relying on the Min Common Supergraph for minimal injection of redundant code

Conclusion

- PSLP improves vectorization coverage compared to the state-of-the-art
- Converts non-isomorphic code into isomorphic by:
 - Relying on the Min Common Supergraph for minimal injection of redundant code
 - Emitting Select instructions to guarantee correctness

Conclusion

- PSLP improves vectorization coverage compared to the state-of-the-art
- Converts non-isomorphic code into isomorphic by:
 - Relying on the Min Common Supergraph for minimal injection of redundant code
 - Emitting Select instructions to guarantee correctness
 - Optimizing away redundant Selects

Conclusion

- PSLP improves vectorization coverage compared to the state-of-the-art
- Converts non-isomorphic code into isomorphic by:
 - Relying on the Min Common Supergraph for minimal injection of redundant code
 - Emitting Select instructions to guarantee correctness
 - Optimizing away redundant Selects
- PSLP performs better compared to SLP on commodity SIMD-capable hardware